

实 验 报 告

Lab4

姓名：朱瑞媛

学号：14307130373

Exercise 1. `i386_init` identifies the file system environment by passing the type `ENV_TYPE_FS` to your environment creation function, `env_create`. Modify `env_create` in `env.c`, so that it gives the file system environment I/O privilege, but never gives that privilege to any other environment. Make sure you can start the file environment without causing a General Protection fault. You should pass the "fs i/o" test in **make grade**.

```
412
413     if (type == ENV_TYPE_FS)
414         env->env_tf.tf_eflags |= FL_IOPL_MASK;
```

在`env_create`函数中修改进程的`eflag`值。

```
make[1]: Leaving directory '/home/zyj/cdb
internal FS tests [fs/test.c]: OK (1.4s)
fs i/o: OK
```

通过测试。

Question Do you have to do anything else to ensure that this I/O privilege setting is saved and restored properly when you subsequently switch from one environment to another? Why?

在进程切换时调用了`env.pop_tf`函数，其中进行了寄存器的恢复，在`iret`指令中恢复了`eip`，`cs`，`eflags`等寄存器。

Exercise 2. Implement the `bc_pgfault` and `flush_block` functions in `fs/bc.c`. `bc_pgfault` is a page fault handler, just like the one you wrote in the previous lab for copy-on-write fork, except that its job is to load pages in from the disk in response to a page fault. When writing this, keep in mind that (1) `addr` may not be aligned to a block boundary and (2) `ide_read` operates in sectors, not blocks.

The `flush_block` function should write a block out to disk *if necessary*.

`flush_block` shouldn't do anything if the block isn't even in the block cache (that is, the page isn't mapped) or if it's not dirty. We will use the VM hardware

to keep track of whether a disk block has been modified since it was last read from or written to disk. To see whether a block needs writing, we can just look to see if the PTE_D "dirty" bit is set in the uvpt entry. (The PTE_D bit is set by the processor in response to a write to that page; see 5.2.4.3 in chapter 5 of the 386 reference manual.) After writing the block to disk, flush_block should clear the PTE_D bit using sys_page_map. Use **make grade** to test your code. Your code should pass "check_bc", "check_super", and "check_bitmap".

```

78 void
79 flush_block(void *addr)
80 {
81     uint32_t blockno = ((uint32_t)addr - DISKMAP) / BLKSIZE;
82
83     if (addr < (void*)DISKMAP || addr >= (void*)(DISKMAP + DISKSIZE))
84         panic("flush_block of bad va %08x", addr);
85
86     // LAB 5: Your code here.
87     // panic("flush_block not implemented");
88     int r;
89     addr = ROUNDDOWN(addr, PGSIZE);
90     if (va_is_mapped(addr) && va_is_dirty(addr))
91     {
92         if ((r = ide_write(blockno * BLKSECTS, addr, BLKSECTS)) < 0)
93             panic("in flush_block, ide_write.\n");
94         if ((r = sys_page_map(0, addr, 0, addr, uvpt[PGNUM(addr)] & PTE_SYSCALL)) < 0)
95             panic("in flush_block, sys_page_map.\n");
96     }
97 }
98

```

先根据地址计算对应的blockno，然后检查正确性，最后判断是否是脏块，如果是则写回磁盘并清除dirty位。

```

43     panic("reading non-existent block %08x\n", blockno);
44
45     // Allocate a page in the disk map region, read the contents
46     // of the block from the disk into that page.
47     // Hint: first round addr to page boundary. fs/ide.c has code to read
48     // the disk.
49     //
50     // LAB 5: you code here:
51
52     addr = ROUNDDOWN(addr, PGSIZE);
53     if ((r = sys_page_alloc(0, addr, PTE_U | PTE_P | PTE_W)) < 0)
54         panic("in bc_pgfault, sys_page_alloc.\n");
55     if ((r = ide_read(blockno * BLKSECTS, addr, BLKSECTS)) < 0)
56         panic("in bc_pgfault, ide_read.\n");
57

```

先根据地址计算出对应的blockno，然后检查正确性包括地址是否在映射范围内、对应的block是否存在等。

```

internal FS tests [fs/test.c]: OK (1.3s)
ashfsigi/o: OK
check_bc: OK
check_super: OK
check_bitmap: OK

```

通过测试。

Exercise 3. Use `free_block` as a model to implement `alloc_block` in `fs/fs.c`, which should find a free disk block in the bitmap, mark it used, and return the number of that block. When you allocate a block, you should immediately flush the changed bitmap block to disk with `flush_block`, to help file system consistency. Use **make grade** to test your code. Your code should now pass "alloc_block".

```

57 int
58 alloc_block(void)
59 {
60     // The bitmap consists of one or more blocks. A single bitmap block
61     // contains the in-use bits for BLKBITSIZE blocks. There are
62     // super->s_nblocks blocks in the disk altogether.
63
64     // LAB 5: Your code here.
65     uint32_t blockno;
66
67     for(blockno = 0; blockno < super->s_nblocks; blockno++)
68     {
69         if (block_is_free(blockno))
70         {
71             bitmap[blockno/32] ^= 1 << (blockno % 32);
72             flush_block(bitmap);
73             return blockno;
74         }
75     }

```

以`free_block`为参考实现`alloc_block`，功能是在位图中查找1个空闲磁盘块，标记为占用并返回块序号。当分配1个块时，为了维护文件系统的一致性，需要快速地使用`flush_block`函数写回你对位图的修改。

```

check_super: OK
check_bitmap: OK
alloc_block: OK

```

通过测试。

Exercise 4. Implement `file_block_walk` and `file_get_block`. `file_block_walk` maps from a block offset within a file to the pointer for that block in the struct `File` or the indirect block, very much like what `pgdir_walk` did for page tables. `file_get_block` goes one step further and maps to the actual disk block, allocating a new one if necessary.

```
148 static int
149 file_block_walk(struct File *f, uint32_t filebno, uint32_t **ppdiskbno, bool alloc)
150 {
151     // LAB 5: Your code here.
152
153     int r;
154
155     if (filebno >= NDIRECT + NINDIRECT)
156         return -E_INVALID;
157     if (filebno < NDIRECT)
158     {
159         if (ppdiskbno)
160             *ppdiskbno = f -> f_direct + filebno;
161         return 0;
162     }
163     if (!alloc && !f -> f_indirect)
164         return -E_NOT_FOUND;
165     if (!f -> f_indirect)
166     {
167         if ((r = alloc_block()) < 0)
168             return -E_NO_DISK;
169         f -> f_indirect = r;
170         memset(diskaddr(r), 0, BLKSIZE);
171         flush_block(diskaddr(r));
172     }
173     if (ppdiskbno)
174         *ppdiskbno = (uint32_t *)diskaddr(f -> f_indirect) + filebno - NDIRECT;
175     return 0;
176
177     //panic("file_block_walk not implemented");
178 }
179 }
```

`file_block_walk`函数寻找一个文件结构`f`中的第`filebno`个块指向的磁盘块编号放入`ppdiskbno`。

```

189 int
190 file_get_block(struct File *f, uint32_t filebno, char **blk)
191 {
192     // LAB 5: Your code here.
193     int r;
194     uint32_t *ppdiskbno;
195
196     if ((r = file_block_walk(f, filebno, &ppdiskbno, 1)) < 0)
197         return r;
198     if (*ppdiskbno == 0)
199     {
200         if ((r = alloc_block()) < 0)
201             return -E_NO_DISK;
202         *ppdiskbno = r;
203         memset(diskaddr(r), 0, BLKSIZE);
204         flush_block(diskaddr(r));
205     }
206
207     *blk = diskaddr(*ppdiskbno);
208     return 0;
209
210     //panic("file_get_block not implemented");
211 }

```

file_get_block函数先调用file_walk_block函数找到文件中的目标块，然后将其转换为地址空间中的地址赋值给blk。

```

check_bc: OK
check_super: OK
check_bitmap: OK
alloc_block: OK
file_open: OK
file_get_block: OK
file_flush/file_truncate/file rewrite: OK
testfile: OK (1.1s)

```

通过测试。

Exercise 5. Implement serve_read in fs/serv.c. serve_read's heavy lifting will be done by the already-implemented file_read in fs/fs.c (which, in turn, is just a bunch of calls to file_get_block). serve_read just has to provide the RPC interface for file reading. Look at the comments and code in serve_set_size to get a general idea of how the server functions should be structured.

```

207 int
208 serve_read(envid_t envid, union Fsipc *ipc)
209 {
210     struct Fsreq_read *req = &ipc->read;
211     struct Fsret_read *ret = &ipc->readRet;
212
213     if (debug)
214         cprintf("serve_read %08x %08x %08x\n", envid, req->req_fileid, req->req_n);
215
216     // Lab 5: Your code here:
217
218     struct OpenFile *o;
219     int r, req_n;
220
221     if ((r = openfile_lookup(envid, req -> req_fileid, &o)) < 0)
222         return r;
223     req_n = req -> req_n > PGSIZE ? PGSIZE : req -> req_n;
224     if ((r = file_read(o -> o_file, ret -> ret_buf, req_n, o -> o_fd -> fd_offset)) < 0)
225         return r;
226     o -> o_fd -> fd_offset += r;
227     return r;
228
229 }

```

先从Fsipc中获取读请求的结构体,然后在openfile中查找fileid对应的Openfile结构体,紧接着从openfile长相的o_file中读取内容到保存返回结果的ret_buf中,并移动文件偏移指针。

```

file_flush/file_truncate/file rewrite: OK
testfile: foke(1.0s) starting at
the serve_open/file_stat/file_close: OK
s arfile_read: so the number of

```

通过测试。

Exercise 6. Implement serve_write in fs/serv.c and devfile_write in lib/file.c.

```

236 int
237 serve_write(envid_t envid, struct Fsreq_write *req)
238 {
239     if (debug)
240         cprintf("serve_write %08x %08x %08x\n", envid, req->req_fileid, req->req_n);
241
242     // LAB 5: Your code here.
243
244     struct OpenFile *o;
245     int r, req_n;
246
247     if ((r = openfile_lookup(envid, req -> req_fileid, &o)) < 0)
248         return r;
249     req_n = req -> req_n > PGSIZE ? PGSIZE : req -> req_n;
250     if ((r = file_write(o -> o_file, req -> req_buf, req_n, o -> o_fd -> fd_offset)) < 0)
251         return r;
252     o -> o_fd -> fd_offset += r;
253     return r;
254
255
256     //panic("serve_write not implemented");
257 }

```

serve_write基本和serve_read是一样的。调用了file_write函数。

```

136 static ssize_t
137 devfile_write(struct Fd *fd, const void *buf, size_t n)
138 {
139     // Make an FSREQ_WRITE request to the file system server. Be
140     // careful: fsipcbuf.write.req_buf is only so large, but
141     // remember that write is always allowed to write *fewer*
142     // bytes than requested.
143     // LAB 5: Your code here
144
145     int r;
146     if (n > sizeof(fsipcbuf.write.req_buf))
147         n = sizeof(fsipcbuf.write.req_buf);
148     fsipcbuf.write.req_fileid = fd -> fd_file.id;
149     fsipcbuf.write.req_n = n;
150     memmove(fsipcbuf.write.req_buf, buf, n);
151     if ((r = fsipc(FSREQ_WRITE, NULL)) < 0)
152         return r;
153     return r;
154
155     // panic("devfile_write not implemented");
156 }

```

devfile_write同样是一个用户库,功能是打包各种参数,然后调用IPC,请求FS内核进程做读取操作。因为是写操作,所以需要 用memmove来把buf里的内容移动到fsipcbuf.write.req_buf.

```

) serve_open/file_stat/file_close: OK
  file_read: OK
  file_write: OK
; file_read after file_write: OK
  open: OK
  large file: OK

```

通过测试。